



Figure 1: logo

## Octopus

Copyright 2013 The Netherlands eScience Center

Author: Jason Maassen ([J.Maassen@esciencecenter.nl](mailto:J.Maassen@esciencecenter.nl))

Version: Userguide v0.2, Octopus v1.0rc1

Last modified: 22 August 2013

## Copyrights & Disclaimers

Octopus is copyrighted by the Netherlands eScience Center and releases under the Apache License, Version 2.0.

See the “LICENSE” and “NOTICE” files in the octopus distribution for more information.

For more information on the Netherlands eScience Center see:

<http://www.esciencecenter.nl>

The octopus project web site can be found at:

<https://github.com/NLeSC/octopus>.

This product includes the SLF4J library, which is Copyright (c) 2004-2013 QOS.ch See “notices/LICENSE.slf4j.txt” for the licence information of the SLF4J library.

This product includes the JSch library, which is Copyright (c) 2002-2012 Atsuhiko Yamanaka, JCraft, Inc. See “notices/LICENSE.jsch.txt” for the licence information of the JSch library.

## What is it?

Octopus is a middleware abstraction library. It provides a simple Java programming interface to various pieces of software that can be used to access distributed compute and storage resources.

## Why Octopus?

Octopus is developed by the Netherlands eScience Center as a support library for our projects. Several projects develop end-user applications that require access to distributed compute and storage resources. Octopus provides a simple API to access those resources, allowing those applications to be developed more rapidly. The experience gained during the development of these end-user applications is used to improve the Octopus API and implementation.

## Installation

The installation procedure and dependencies of the octopus library can be found in the file “INSTALL.md” in the octopus distribution.

## Design

Octopus is designed with extensibility in mind. It uses a modular and layer design as shown in Figure 1.

Octopus consists of three layers, an *interface layer*, an *engine layer* and an *adaptor layer*.

The *interface layer* is used by the application using octopus. It contains several specialized interfaces:

- Octopus: this is the main entry point used to retrieve the other interfaces.
- Files: contains functionality related to files, e.g., creation, deletion, copying, reading, writing, obtaining directory listings, etc.
- Jobs: contains functionality related to job submission, e.g., submitting, polling status, cancelling, etc.
- Credentials: contains functionality related to credentials. Credentials (such as a username password combination) are often needed to gain access to files or to submit jobs.

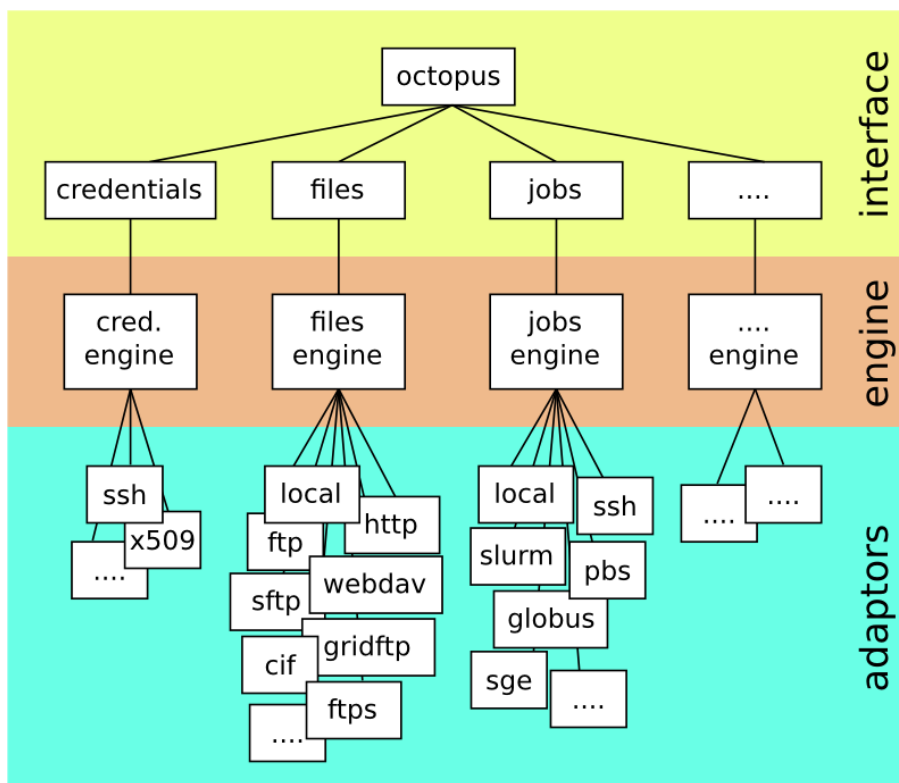


Figure 2: Octopus design

The modular design of octopus allows us to add additional interfaces in later versions, e.g., a Clouds interface to manage virtual machines, or a Networks interface to manage bandwidth-on-demand networks.

The *adaptor layer* contains the adaptors for the each of the middlewares that octopus supports. An *adaptor* offers a middleware specific implementation for the functionality offered by one of the interfaces in octopus.

For example, an adaptor may provide an *sftp* specific implementation of the functions in the octopus *file interface* (such as *copy* or *delete*) by translating each of these functions to *sftp* specific code and commands.

For each interface in octopus there may be multiple adaptors translating its functionality to different middlewares. To distinguish between these adaptors, octopus uses the *scheme* they support, such as “sftp”, “http” or “ssh”. There can be only one adaptor for each scheme.

The *engine layer* of octopus contains the “glue” that connects each interface to the adaptors that implement its functionality. When a function of the interface layer is invoked, the call will be forwarded to the engine layer. It is then the responsibility of the engine layer to forward this call to the right adaptor.

To perform this selection, the engine layer matches the *scheme* of the object on which the operation needs to be performed, to the *schemes* supported by each of the adaptors. When the schemes match, the adaptor is selected.

## Interfaces and datatypes

This section will briefly explain each of the interfaces and related datatypes. Detailed information about Octopus can be found in the online JavaDOC:

<http://nlesc.github.io/octopus/javadoc/>

### Package Structure

The octopus API uses the following package structure:

- `nl.esciencecenter.octopus` Entry point into octopus.
- `nl.esciencecenter.octopus.credentials` Credential interface.
- `nl.esciencecenter.octopus.files` Files interface.
- `nl.esciencecenter.octopus.jobs` Jobs interface.
- `nl.esciencecenter.octopus.exceptions` Exceptions used in octopus.
- `nl.esciencecenter.octopus.util` Various utility classes (experimental).

We will now briefly describe the most important classes and interfaces of these packages.

## Octopus factory and interface

The `nl.esciencecenter.octopus` contains the entry point into the octopus library. It contains the `OctopusFactory` class and `Octopus` interface.

```
public class OctopusFactory {
    public static Octopus newOctopus(Map<String,String> properties) throws ...
    public static void endOctopus(Octopus octopus) throws ...
    public static void endAll();
}
```

The `OctopusFactory` class contains the `newOctopus` method to create a new octopus instance. The `endOctopus` method can be used to end the octopus instance once it is no longer needed. It is important to end the octopus when it is no longer needed, as this allows it to release any resources it has obtained.

When creating an octopus using `newOctopus`, the `properties` parameter can be used to configure the octopus instance. If no configuration is necessary, `null` can be used. Properties consist of a set of key-value pairs. In octopus all keys **must** start with "octopus.". To configure the adaptors, properties of the form "octopus.adaptors.<name>.<property>" can be used, where <name> is the name of the adaptor (for example `local` or `ssh`) and <property> is the name of the property to be configured. Note that this name can be further qualified, for example `octopus.adaptors.local.a.b.c`. The available properties can be found in the documentation of the individual adaptors (see Appendix A).

```
public interface Octopus {
    Files files();
    Jobs jobs();
    Credentials credentials();
    Map<String,String> getProperties();
    AdaptorStatus getAdaptorInfo(String adaptorName) throws ...
    AdaptorStatus[] getAdaptorInfos();
}
```

Once an `Octopus` is created using the `newOctopus` method, the `files`, `jobs` and `credentials` methods in this interface can be used to retrieve various interfaces that the octopus library offers. They will be described in more detail below.

The `getProperties` method can be used to retrieve the properties used when the octopus was created. Most objects created by octopus contain such a `getSupportedProperties` method. For brevity, we will not explain these further.

The `getAdaptorInfo` and `getAdaptorInfos` methods can be used to retrieve information about the adaptors. This information is returned in an `AdaptorStatus` object:

```

public interface AdaptorStatus {
    String getName();
    String getDescription();
    String[] getSupportedSchemes();
    OctopusPropertyDescription[] getSupportedProperties();
    Map<String, String> getAdaptorSpecificInformation();
}

```

An `AdaptorStatus` contains methods to retrieve the name of an adaptor (`getName`), get a (human readable) description of what functionality it has to offer (`getDescription`) and retrieve a list of the schemes it supports (`getSupportedSchemes`).

The `getSupportedProperties` can be used to retrieve a list of configuration options the adaptor supports. Each `OctopusPropertyDescription` gives a full description of a single property, including its name (of the form "octopus.adaptors.<name>.<property>"), the expected type of its value, a human readable description of its purpose, etc. More information can be found in Appendix A.

Finally, `getAdaptorSpecificInformation` can be used to retrieve status information from the adaptor. Each key contains a property of the form described above. The possible returned properties can be found in the *Adaptor* section below.

## Credentials interface

The `nl.esciencecenter.octopus.credentials` package contains the credentials interface of octopus. The main entrypoint is `Credentials`:

```

public interface Credentials {

    Credential newCertificateCredential(String scheme, String keyfile,
        String certfile, String username, char [] password,
        Map<String,String> properties) throws ...;

    Credential newPasswordCredential(String scheme, String username,
        char [] password, Map<String,String> properties) throws ...;

    Credential getDefaultCredential(String scheme) throws ...;
}

```

The `Credentials` interface contains various methods for creating credentials, based on certificates or passwords. For each method, the desired *scheme* needs to be provided as a parameter. This allows octopus to forward the call to the correct

adaptor. Note that some types of credentials may not be supported by all adaptors. An exception will be thrown when an unsupported `new<Type>Credential` methods is invoked.

Additional configuration can also be provides using the `properties` parameter, which use the same form as described in the *Octopus factory and interface* section above. If no additional configuration is needed, `null` can be used. The `getDefaultCredential` returns the default credential for the given scheme. All adaptors are guarenteed to support this method.

All `new<Type>Credential` methods return a `Credential` that contains the following methods:

```
public interface Credential {
    String getAdaptorName();
    Map<String,String> getProperties();
}
```

The `getAdaptorName` method can be used to retrieve the name of the adaptor that created the credential. Many adaptor specific objects returned by octopus contain this method. For brevity we will not explain this further.

## Files interface

The `nl.esciencecenter.octopus.files` package contains the files interface of octopus. The main entrypoint is `Files`. For readability we will split the explanation of `Files` into several parts:

```
public interface Files {

    FileSystem newFileSystem(Uri location, Credential credential,
        Map<String,String> properties) throws ...

    FileSystem getLocalCWDFileSystem() throws ...

    FileSystem getLocalHomeFileSystem() throws ...

    void close(FileSystem filesystem) throws ...

    boolean isOpen(FileSystem filesystem) throws ...

    // ... more follows

}
```

The `Files` interface contains several methods for creating and closing a `FileSystem`. A `FileSystem` provides an abstraction for a (possibly remote) file system. To create a `FileSystem` the `newFileSystem` method can be used. The `URI location` parameter provides the information on the location of the file system. The `URI` is expected to contain at least a *scheme*. Most `URIs` will also contain *host* information. Optionally, *user* information may also be provided. A file system `URI` may *not* contain a path other than `"/"`. The following are all valid file system `URIs`:

```
file:///
sftp://example.com
sftp://test@example.com:8080/
```

The `newFileSystem` method also has a `credential` parameter to provide the credential needed to access the file system. If this parameter is set to `null` the default credentials will be used for the scheme. The `properties` parameter can be used to provide additional configuration properties. Again, `null` can be used if no additional configuration is required. The returned `FileSystem` contains the following:

```
public interface FileSystem {
    /// ...
    URI getUri();
    AbsolutePath getEntryPath();
}
```

The `getUri` returns the `URI` used to create it. The `getEntryPath` method returns the *path at which the file system was entered*. For example, when accessing a file system using `sftp` it is customary (but not mandatory) to enter the file system at the users' home directory. Therefore, the entry path of the `FileSystem` will be `/home/username`.

The `getLocalCWDFileSystem` and `getLocalHomeFileSystem` methods of `Files` provide shortcuts to create a `FileSystem` representing the *current working directory* or *user home directory* on the local machine.

When a `FileSystem` is no longer used, it **must** be closed using `close`. This releases any resources held by the `FileSystem`. The `isOpen` method can be used to check if a `FileSystem` is open or closed.

Once a `FileSystem` is created, it can be used to access files:

```
public interface Files {

    AbsolutePath newPath(FileSystem filesystem,
        Pathname location) throws ...
}
```



```

    Path createFile(Path path) throws ...

    Path createDirectories(Path dir) throws ...

    Path createDirectory(Path dir) throws ...

    boolean exists(Path path) throws ...

    void delete(Path path) throws ...

    FileAttributes getAttributes(Path path) throws ...

    // ... more follows
}

```

The `newPath` method can be used to create a new `Path`. An `Path` represents a path on a specific `FileSystem`. This path does not necessarily exist. To create an `Path`, both the target `FileSystem` and a `Pathname` are needed. A `Pathname` contains a sequence of strings separated using a special *separator* character, which is used to identify a location on a file system. For example `/tmp/dir` or `c:\windows\Users`. `Pathname` contains many utility methods for manipulating these string sequences. The details can be found in the Javadoc.

`Files` contains several methods to create and delete files and directories. When creating files and directories octopus checks if the target already exists. If so, an exception will be thrown. Similarly, an exception is thrown when attempting to delete non-existing file or a directory that is not empty. The `exists` method can be used to check if a path exists.

Using the `getAttributes` method the attributes of a file can be retrieved. These `FileAttributes` contain information on the type of file (regular file, directory, link, etc), its size, creation time, access rights, etc.

To list directories, the following methods are available:

```

public interface Files {

    DirectoryStream<Path> newDirectoryStream(Path dir) throws ...

    DirectoryStream<PathAttributesPair> newAttributesDirectoryStream(Path dir) throws ...

    // ... more follows
}

```

Both `newDirectoryStream` and `newAttributesDirectoryStream` return a `DirectoryStream` which can be used to iterate over the contents of a directory.

For the latter, the `FileAttributes` for each of the files are also included. alternatively, these methods are also available with an extra `filter` parameter, which can be used to filter the stream in advance.

To read or write files, the following methods are available:

```
public interface Files {

    InputStream newInputStream(Path path) throws ...

    OutputStream newOutputStream(Path path, OpenOption... options) throws ...
}
```

Using these methods, an `InputStream` can be created to read a file, and an `OutputStream` can be created to write a file. The `newOutputStream` method requires a `OpenOption... options` parameter to specify how the file should be opened for writing (for example, should the data be append or should the file be truncated first). These options are describe in more detail in the Javadoc.

To copy files, the following methods are available:

```
public interface Files {

    Copy copy(Path source, Path target, CopyOption... options) throws ...

    CopyStatus getCopyStatus(Copy copy) throws ...

    CopyStatus cancelCopy(Copy copy) throws ...

}
```

The `copy` method supports various copy operations such as a regular copy, a resume or an append. The `CopyOption...options` parameter can be used to specify the desired operation. The details can be found in the Javadoc.

Normally, `copy` performs its operation *synchronously*, that is, the call blocks until the copy is completed. However, *asynchronous* operations are also supported by providing the option `CopyOption.ASYNCHRONOUS`. In that case a `Copy` object is returned that can be used to retrieve the status of the copy (using `getCopyStatus`) or cancel it (using `cancelCopy`).

## Jobs interface

The `nl.esciencecenter.octopus.job` package contains the job interface of octopus. The main entrypoint is `Jobs`. For readability we will split the explanation of `Jobs` into several parts:

```

public interface Jobs {

    Scheduler newScheduler(Uri location, Credential credential,
        Map<String,String> properties) throws ...

    Scheduler getLocalScheduler() throws ...
    void close(Scheduler scheduler) throws ...
    boolean isOpen(Scheduler scheduler) throws ...

    // ... more follows
}

```

The `Jobs` interface contains two methods to create a `Scheduler`. A `Scheduler` provides an abstraction for a (possibly remote) scheduler that can be used to run jobs. To create a new scheduler, the `newScheduler` method can be used, which, similar to `newFileSystem`, has `Uri`, `Credential` and `Properties` as parameters. For an explanation of these parameters see `newFileSystem`. `Jobs` also contains a shortcut method `getLocalScheduler` to create a new `Scheduler` for the local machine.

When a `Scheduler` is no longer used, it **must** be closed using the `close` method. The `isOpen` method can be used to check if a `Scheduler` is open or closed.

A `Scheduler` contains the following:

```

public interface Scheduler {

    String[] getQueueNames();
    boolean isOnline();
    boolean supportsInteractive();
    boolean supportsBatch();

    // ...
}

```

Each `Scheduler` contains one or more queues to which jobs can be submitted. Each queue has a name that is unique to the `Scheduler`. The `getQueueNames` method can be used to retrieve all queue names.

The `isOnline` method can be used to determine if the `Scheduler` is an *online scheduler* or an *offline scheduler*. Online schedulers need to remain active for their jobs to run. Ending an online scheduler will kill all jobs that were submitted to it. Offline schedulers do not need to remain active for their jobs to run. A submitted job will typically be handed over to some external server that will manage the job for the rest of its lifetime.

The `supportsInteractive` and `supportsBatch` method can be use to check if the `Scheduler` supports interactive and/or batch jobs. This will be explained below.

Once a `Scheduler` is created, `Jobs` contains several methods to retrieve information about the `Scheduler`.

```
public interface Jobs {

    String getDefaultQueueName(Scheduler scheduler) throws ...

    QueueStatus getQueueStatus(Scheduler scheduler,
        String queueName) throws ...

    QueueStatus[] getQueueStatuses(Scheduler scheduler,
        String... queueNames) throws ...

    Job[] getJobs(Scheduler scheduler,
        String... queueNames) throws ...

    // ... more follows
}
```

The `getQueueStatuses` method can be used to retrieve information about a queue. If no queue names are provided as a parameter, information on all queues in the scheduler will be returned. Using the `getDefaultQueueName` the default queue can be retrieved for the `Scheduler`. The `getJobs` method can be used to retrieve information on all jobs in a queue. Note that this may also include jobs from other users.

To submit and manage jobs, the `Jobs` interface contains the following methods:

```
public interface Jobs {

    Job submitJob(Scheduler scheduler,
        JobDescription description) throws ...

    Streams getStreams(Job job) throws ...

    JobStatus getJobStatus(Job job) throws ...

    JobStatus[] getJobStatuses(Job... jobs);

    JobStatus waitUntilRunning(Job job, long timeout) throws ...

    JobStatus waitUntilDone(Job job, long timeout) throws ...
}
```

```

        JobStatus cancelJob(Job job) throws ...
    }

```

The `submitJob` method can be used to submit a job to a `Scheduler`. A `JobDescription` must be provided as parameter. A `JobDescription` contains all necessary information on how to start the job, for example, the location of the executable, any command line arguments that are required, the working directory, etc. See the Javadoc for details of the `JobDescription`.

Once a job is submitted, a `Job` object is returned that can be used later to retrieve the status of the job (`getJobStatus` or `getJobStatuses`) or to cancel it (`cancelJob`). This `Job` contains the following:

```

public interface Job {
    JobDescription getJobDescription();
    Scheduler getScheduler();
    String getIdentifier();
    boolean isInteractive();
    boolean isOnline();
}

```

Besides methods for retrieving the `JobDescription` and `Scheduler` that created it, each `Job` also contains methods to determine if the `Job` is running on an online `Scheduler` (`isOnline`) and whether the `Job` is an interactive or batch job (`isInteractive`).

Interactive jobs are jobs where the user gets direct control over the standard streams of the job (the *stdin*, *stdout* and *stderr* streams). The user **must** retrieve these streams using the `getStreams` method in `Jobs` and then provide input and output, or close the streams. Failing to do so may cause the job to block indefinitely.

Batch jobs are jobs where the standard streams are redirected from and to files. The source and targets for this redirection can be set in the `JobDescription`. See the Javadoc of `JobDescription` for details.

After submitting a job, `waitUntilRunning` can be used to wait until a job is no longer waiting in the queue and `waitUntilDone` can be used to wait until the job has finished.

For all methods returning a `JobStatus`, the following rule applies: after a job has finished, the status is only guaranteed to be returned *once*. Any subsequent calls to a method that returns a `JobStatus` *may* throw an exception stating that the job does not exist. Some adaptors may return a result however.

## Exceptions

The `nl.esciencecenter.octopus.exceptions` package contains the exceptions that may be thrown by octopus. See the Javadoc for the available exceptions.

## Utilities classes

The `nl.esciencecenter.octopus.util` package contains various utility classes. **This package is experimental and not yet ready for use!!**

## Examples

Many examples of how to use octopus can be found online. They will be listed here in order of increasing complexity:

### Initializing Octopus

<https://github.com/NLeSC/octopus/FIXME>

<https://github.com/NLeSC/octopus/FIXME>

### Creating Credentials

<https://github.com/NLeSC/octopus/FIXME>

<https://github.com/NLeSC/octopus/FIXME>

### File Access

<https://github.com/NLeSC/octopus/FIXME>

<https://github.com/NLeSC/octopus/FIXME>

### Job Submission

<https://github.com/NLeSC/octopus/FIXME>

<https://github.com/NLeSC/octopus/FIXME>

## Appendix A: Adaptor Documentation

This section contains the adaptor documentation which is generated from the information provided by the adaptors themselves.

Octopus currently supports 4 adaptors: local, ssh, gridengine, slurm.

## **Adaptor: local**

The local adaptor implements all functionality with standard java classes such as `java.lang.Process` and `java.nio.file.Files`.

### **Supported schemes:**

local, file

### **Supported properties:**

`octopus.adaptors.local.queue.pollingDelay`

The polling delay for monitoring running jobs (in milliseconds).

- Expected type: INTEGER
- Default value: 1000
- Valid for: [OCTOPUS]

`octopus.adaptors.local.queue.multi.maxConcurrentJobs`

The maximum number of concurrent jobs in the multiq..

- Expected type: INTEGER
- Default value: 4
- Valid for: [OCTOPUS]

## **Adaptor: ssh**

The SSH adaptor implements all functionality with remove ssh servers.

### **Supported schemes:**

ssh, sftp

### Supported properties:

#### `octopus.adaptors.ssh.autoAddHostKey`

Automatically add unknown host keys to `known_hosts`.

- Expected type: BOOLEAN
- Default value: true
- Valid for: [FILESYSTEM, SCHEDULER]

#### `octopus.adaptors.ssh.strictHostKeyChecking`

Enable strict host key checking.

- Expected type: BOOLEAN
- Default value: true
- Valid for: [FILESYSTEM, SCHEDULER]

#### `octopus.adaptors.ssh.loadKnownHosts`

Load the standard `known_hosts` file.

- Expected type: BOOLEAN
- Default value: true
- Valid for: [OCTOPUS]

#### `octopus.adaptors.ssh.queue.pollingDelay`

The polling delay for monitoring running jobs (in milliseconds).

- Expected type: LONG
- Default value: 1000
- Valid for: [SCHEDULER]



**octopus.adaptors.ssh.queue.multi.maxConcurrentJobs**

The maximum number of concurrent jobs in the multiq..

- Expected type: INTEGER
- Default value: 4
- Valid for: [SCHEDULER]

**octopus.adaptors.ssh.gateway**

The gateway machine used to create an SSH tunnel to the target.

- Expected type: STRING
- Default value: null
- Valid for: [FILESYSTEM, SCHEDULER]

## **Adaptor: gridengine**

The SGE Adaptor submits jobs to a (Sun/Oracle/Univa) Grid Engine scheduler. This adaptor uses either the local or the ssh adaptor to gain access to the scheduler machine.

### **Supported schemes:**

ge, sge

### **Supported properties:**

**octopus.adaptors.gridengine.ignore.version**

Skip version check is skipped when connecting to remote machines. WARNING: it is not recommended to use this setting in production environments!

- Expected type: BOOLEAN
- Default value: false
- Valid for: [SCHEDULER]

**octopus.adaptors.gridengine.accounting.grace.time**

Number of milliseconds a job is allowed to take going from the queue to the qacct output.

- Expected type: LONG
- Default value: 60000
- Valid for: [SCHEDULER]

**octopus.adaptors.gridengine.poll.delay**

Number of milliseconds between polling the status of a job.

- Expected type: LONG
- Default value: 1000
- Valid for: [SCHEDULER]

## **Adaptor: slurm**

The Slurm Adaptor submits jobs to a Slurm scheduler. This adaptor uses either the local or the ssh adaptor to gain access to the scheduler machine.

### **Supported schemes:**

slurm

### **Supported properties:**

**octopus.adaptors.slurm.ignore.version**

Skip version check is skipped when connecting to remote machines. WARNING: it is not recommended to use this setting in production environments!

- Expected type: BOOLEAN
- Default value: false
- Valid for: [SCHEDULER]

**octopus.adaptors.slurm.disable.accounting.usage**

Do not used accounting info of slurm, even when available. Mostly for testing purposes

- Expected type: BOOLEAN
- Default value: false
- Valid for: [SCHEDULER]

**octopus.adaptors.slurm.poll.delay**

Number of milliseconds between polling the status of a job.

- Expected type: LONG
- Default value: 1000
- Valid for: [SCHEDULER]